

ConfD memory usage and performance improvements

Johan Bevemyr


Schema Memory Usage

- Data models are growing
- A single device may have ~700 different models
- Models are rarely used, but must be present on device
- Different model “sets”

FXS files

- An FXS file is primarily a sequence of #cs{} records (ConfSpec), stored in chunks.
- Loading FXSs consists of
 - Reading #cs{} records from files and storing them in ETS tables, which are generic in-memory hash tables.
 - Resolving augmentations between name spaces.
 - An ETS table consists of records for a namespace, and all augmentation into that namespace.
 - In addition to #cs{} records there are #docs{}, #exs_type{}, and header information.

#cs{} record layout

```
-record(cs, {  
    tagpath,  
    htagpath,  
    namespace,  
    hnamespace,  
    exs,   
    keys = [],  
    flags = ?F_CS_READ bor ?F_CS_WRITE,  
    dbm,  
    dba = [],  
    validatemfas = [],  
    actions = [],  
    cmp = 0,  
    hooks = [],  
    hidden = none,  
    notifs = [],  
    symlink = undefined,  
    extra = [],  
    default_ref,  
    secondary_indices = [],  
    cli_flags = 0  
}).
```

```
-record(exs, {  
    tagpath,  
    type,  
    primitive_type,  
    default,  
    attrs = [],  
    min_occurs = 1,  
    max_occurs = 1,  
    children = [],  
    flags = 0  
}).
```

A typical distribution

- revision 2018-08-10 of <http://cisco.com/ns/yang/Cisco-IOS-XE-native>

```
cs: 88.49 MiB
doc: 601.12 KiB
exs_type: 155.57 KiB
fxs_header: 1.34 KiB
info: 896 bytes
augments_header: 40 bytes
```

```
#cs{
  tagpath =
    [address,'ipv6-addrees','eid-cont','database-mapping',ipv6,
      service,'instance-list','instance-container',
      ['http://cisco.com/ns/yang/Cisco-IOS-XE-lisp'|lisp],
      router,native],
  htagpath =
    [1266954393,1103441164,1502191743,1365224135,1228132394,
      855380710,2088542428,608613614,
      [247644804|1303297170],
      532320551,472211213],
  namespace = 'http://cisco.com/ns/yang/Cisco-IOS-XE-native',
  hnamespace = 1270643900,
  exs =
    {exs,
      [address,'ipv6-addrees','eid-cont','database-mapping',ipv6,
        service,'instance-list','instance-container',
        ['http://cisco.com/ns/yang/Cisco-IOS-XE-lisp'|lisp],
        router,native],
      {'urn:ietf:params:xml:ns:yang:ietf-inet-types',
        'ipv6-address'},
      inetAddressIPv6,undefined,[],1,1,[],0},
  keys = [],flags = 618970019642690137449580070,dbm = cdb,
  dba = [],validatemfas = [],actions = [],cmp = 1,hooks = [],
  hidden = none,notifs = [],symlink = undefined,extra = [],
  default_ref = undefined,secondary_indices = [],
  cli_flags = 524288}]
```

Compaction

- Remove redundant information
 - Multiple copies of tagpath
 - htagpath and hnamespace are computed
 - Default values in record

```
-record(cs_trim, {  
    key,  
    dmap,  
    record  
}).
```
- Performance penalty on sync & show ~1-2% slowdown.

Compaction (cont.)


```
#cs_trim{
  key =
    ['Cellular','update-source',neighbor,
    ['http://cisco.com/ns/yang/Cisco-IOS-XE-bgp'|bgp],
    router,native],
  dmap = 229264,
  record =
    {'http://cisco.com/ns/yang/Cisco-IOS-XE-native',
    {exs_trim,1970,
    {'http://www.w3.org/2001/XMLSchema',string},
    string,0},
    618970019642690345755493894,cdb,
    [{cli_allow_join_with_value_r,154742504910672534362390528,
    []}],
    20769187434139310514121985316880384}}
```

Version	ETS table size	ETS table size (compressed)	Ratio
ConfD 6.7	93.82 MiB	56.03 MiB	1
Diff-dep bugfix	82.90 MiB	50.31 MiB	1.1
#cs_trim{}	36.70 MiB	23.54 MiB	2.4

Lazy loading

- Flush all #cs_trim{} records to disk, and only store
 - Path
 - Position in file
- Read records on demand and store in ETS table.

```
-record(cs_file, {  
    tagpath,  
    pos  
}).
```




Lazy loading (cont.)

Version	ETS table size	ETS table size (compressed)	Ratio
ConfD 6.7	93.82 MiB	56.03 MiB	1
Diff-dep bugfix	82.90 MiB	50.31 MiB	1.1
#cs_trim{}	36.70 MiB	23.54 MiB	2.4
#cs_file{}	15.10 MiB	15.79 MiB	3.5

- Pros
 - Use less memory
- Cons
 - Require a baking phase
 - Slower access
 - Somewhat unpredictable memory usage
 - Require disk space

De-duplicate

- Only 10% of #cs{} records are unique when removing tagpath
 - Store #cs_trim{} records under a hash of record
 - Store #cs_ref{} records mapping tagpath to #cs_trim{}



Version	ETS table size	ETS table size (compressed)	Ratio
ConfD 6.7	93.82 MiB	56.03 MiB	1
Diff-dep bugfix	82.90 MiB	50.31 MiB	1.1
#cs_trim{}	36.70 MiB	23.54 MiB	2.4
#cs_file{}	15.10 MiB	15.79 MiB	3.5
#cs_ref{}	19.68 MiB	17.73 MiB	3.2

Path Sharing

- Large parts of data is now tagpaths
- Tag paths are paths from root

```
[]
```

```
[native]
```

```
[native, router]
```

```
[native, router, ['http://cisco.com/ns/yang/Cisco-IOS-XE-bgp'|bgp]]
```

```
[native, router, ['http://cisco.com/ns/yang/Cisco-IOS-XE-bgp'|bgp], neighbor]
```

```
[native, router, ['http://cisco.com/ns/yang/Cisco-IOS-XE-bgp'|bgp], neighbor, 'update-source']
```

```
[native, router, ['http://cisco.com/ns/yang/Cisco-IOS-XE-bgp'|bgp], neighbor, 'update-source', router, nati
```

- Only tail of path is unique, rest is shared with parent

Implementation for Path Sharing

- Custom data structure, mix of trees and hash tables

[d,c,b,a]

[c,b,a]

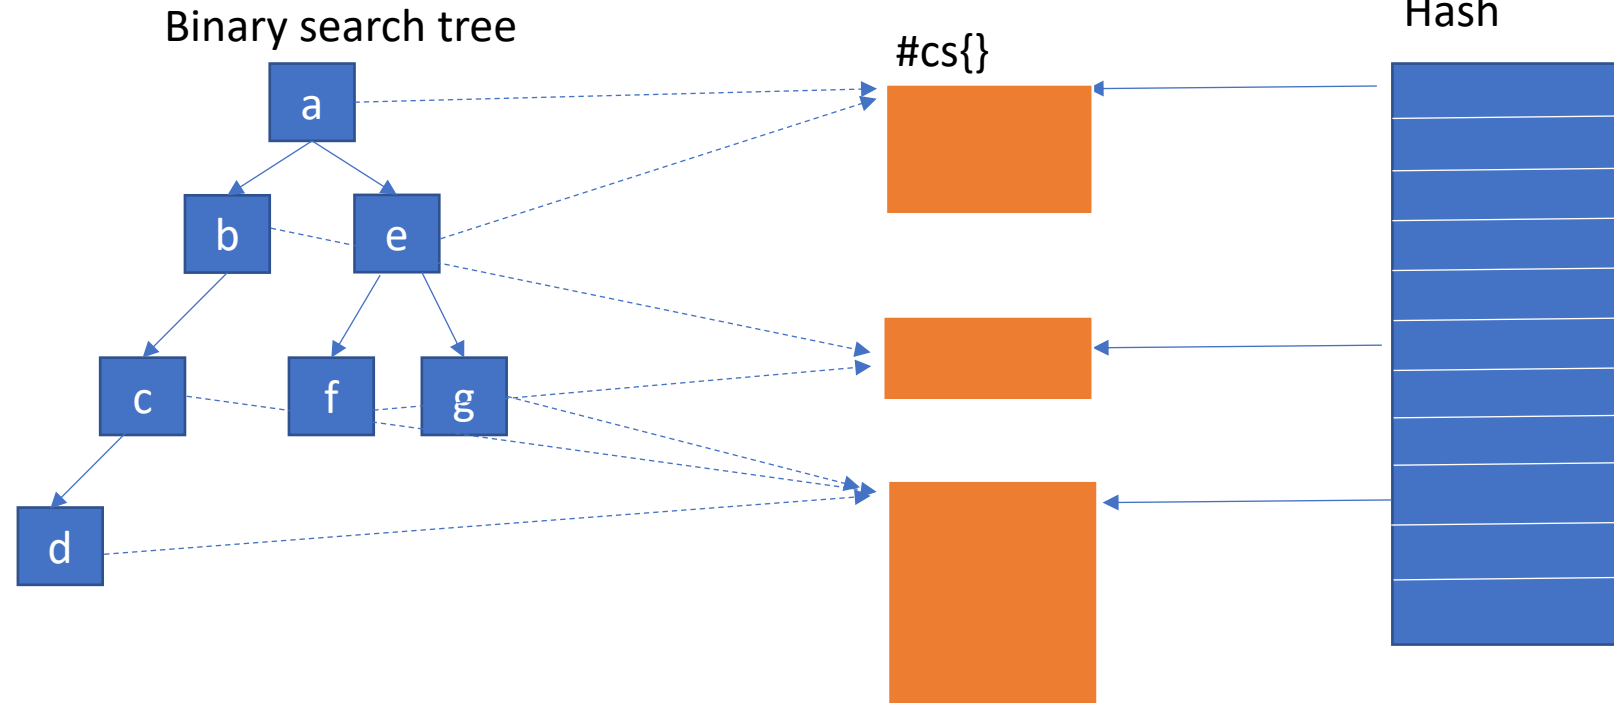
[b,a]

[a]

[f,e,a]

[e,a]

[b,e,a]



New FXS Implementation

Version	ETS table size	ETS table size (compressed)	Ratio
ConfD 6.7	93.82 MiB	56.03 MiB	1
Diff-dep bugfix	82.90 MiB	50.31 MiB	1.1
#cs_trim{}	36.70 MiB	23.54 MiB	2.4
#cs_file{}	15.10 MiB	15.79 MiB	3.5
De-duplicate	19.68 MiB	17.73 MiB	3.2
NIF	9.21 MiB		
NIF trim	8.03 MiB	5.81 MiB	9.6

Approximately twice as fast as the original implementation, and 1/10:th the size. Present in ConfD 7.1.

Cisco-style CLI Parsing

- Large configurations takes a long time to process
- Reading from CLI files (C-style)
- Loading configurations in ConfD
- Native implementation
 - Reduced memory usage compared to Java implementation

CLI: Top-down, breadth-first (on drop-node-name)

Goal: User friendly CLI, provide help and alternatives, precise error reporting

```
interface Loopback0
  ip address 127.0.0.1 255.255.255.255
exit
```

Algorithm:

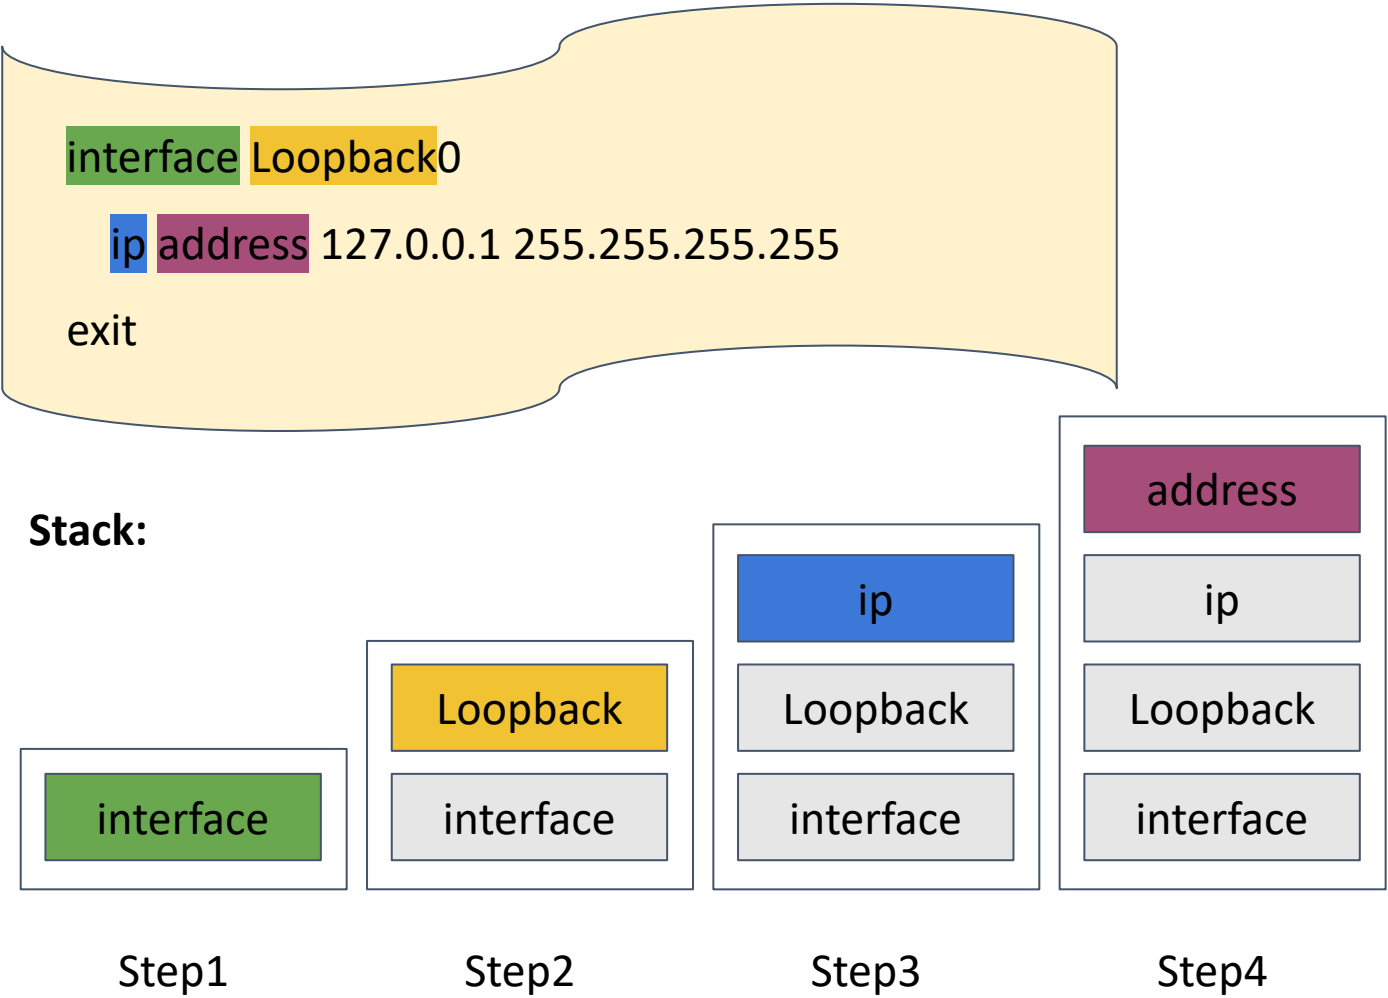
Step1: /: collect interface, policy, mpls, ...	select interface
Step2: /interface: collect Ethernet, ..., Loopback, ...	select Loopback
Step3: /interface/Loopback/: collect ip, ntp, peer, vrf, ...	select ip
Step4: /interface/Loopback/ip: collect access-group, address, arp, rip, ...	select address

YANG tree:

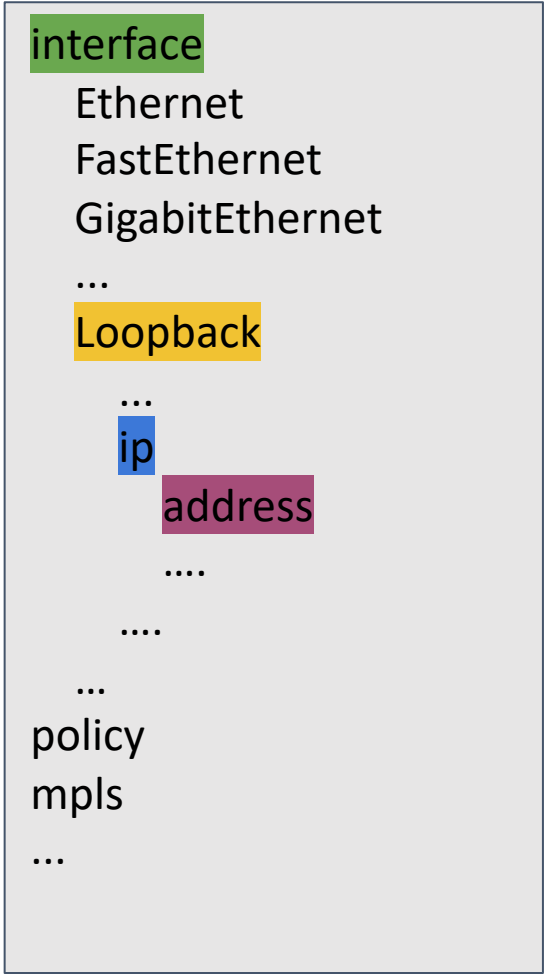
```
interface
  Ethernet
  FastEthernet
  GigabitEthernet
  ...
  Loopback
    ...
    ip
      address
      ....
    ....
  ...
  policy
  mpls
  ...
```


Turbo CLI: Top-down, depth-first, stack based parser

Goal: Fast parsing for programmatic users, pick first best match



YANG tree:



Turbo Parser vs Regular Parser

File	Regular CLI	Turbo Parser	Ratio
1000-lb-config	30 s	5.8 s	5.2
10,000-ace-config-cli-sequence	320 s	26 s	12.3
30,000-ace-config-cli-sequence	940 s	76 s	12.3
huge/nonwireless_CLI.txt	50 s	4.2 s	11.9
huge/ncs4216-all_config.txt	2.7 s	0.67 s	4.0

Filtering on List Instances in Data Provider

Would be nice to pass “filters” to the DP instead of filtering in ConfD – subtree filter, XPath, CLI.

- Avoid round trips between ConfD and data provider
- Simple XPath evaluation in DP

Important when

- Reading large sets of operational data and configuration
- Configuration validation
- Evaluating XPath expressions (e.g. must & when expressions)

Filter Example

```
list foo {  
    key name;  
    leaf name { type string; }  
    leaf value { type int32; }  
}  
  
list bar {  
    key name;  
    leaf name { type string; }  
    leaf fooref {  
        type leafref {  
            path "/foo/value";  
        }  
    }  
}
```

XPath example:

```
/foo[value > 42 or starts-with(name, "eth-")]
```

NETCONF subtree example:

```
<foo>  
  <value>42</value>  
</foo>
```

CLI

```
show foo eth-*  
show foo * value 42
```

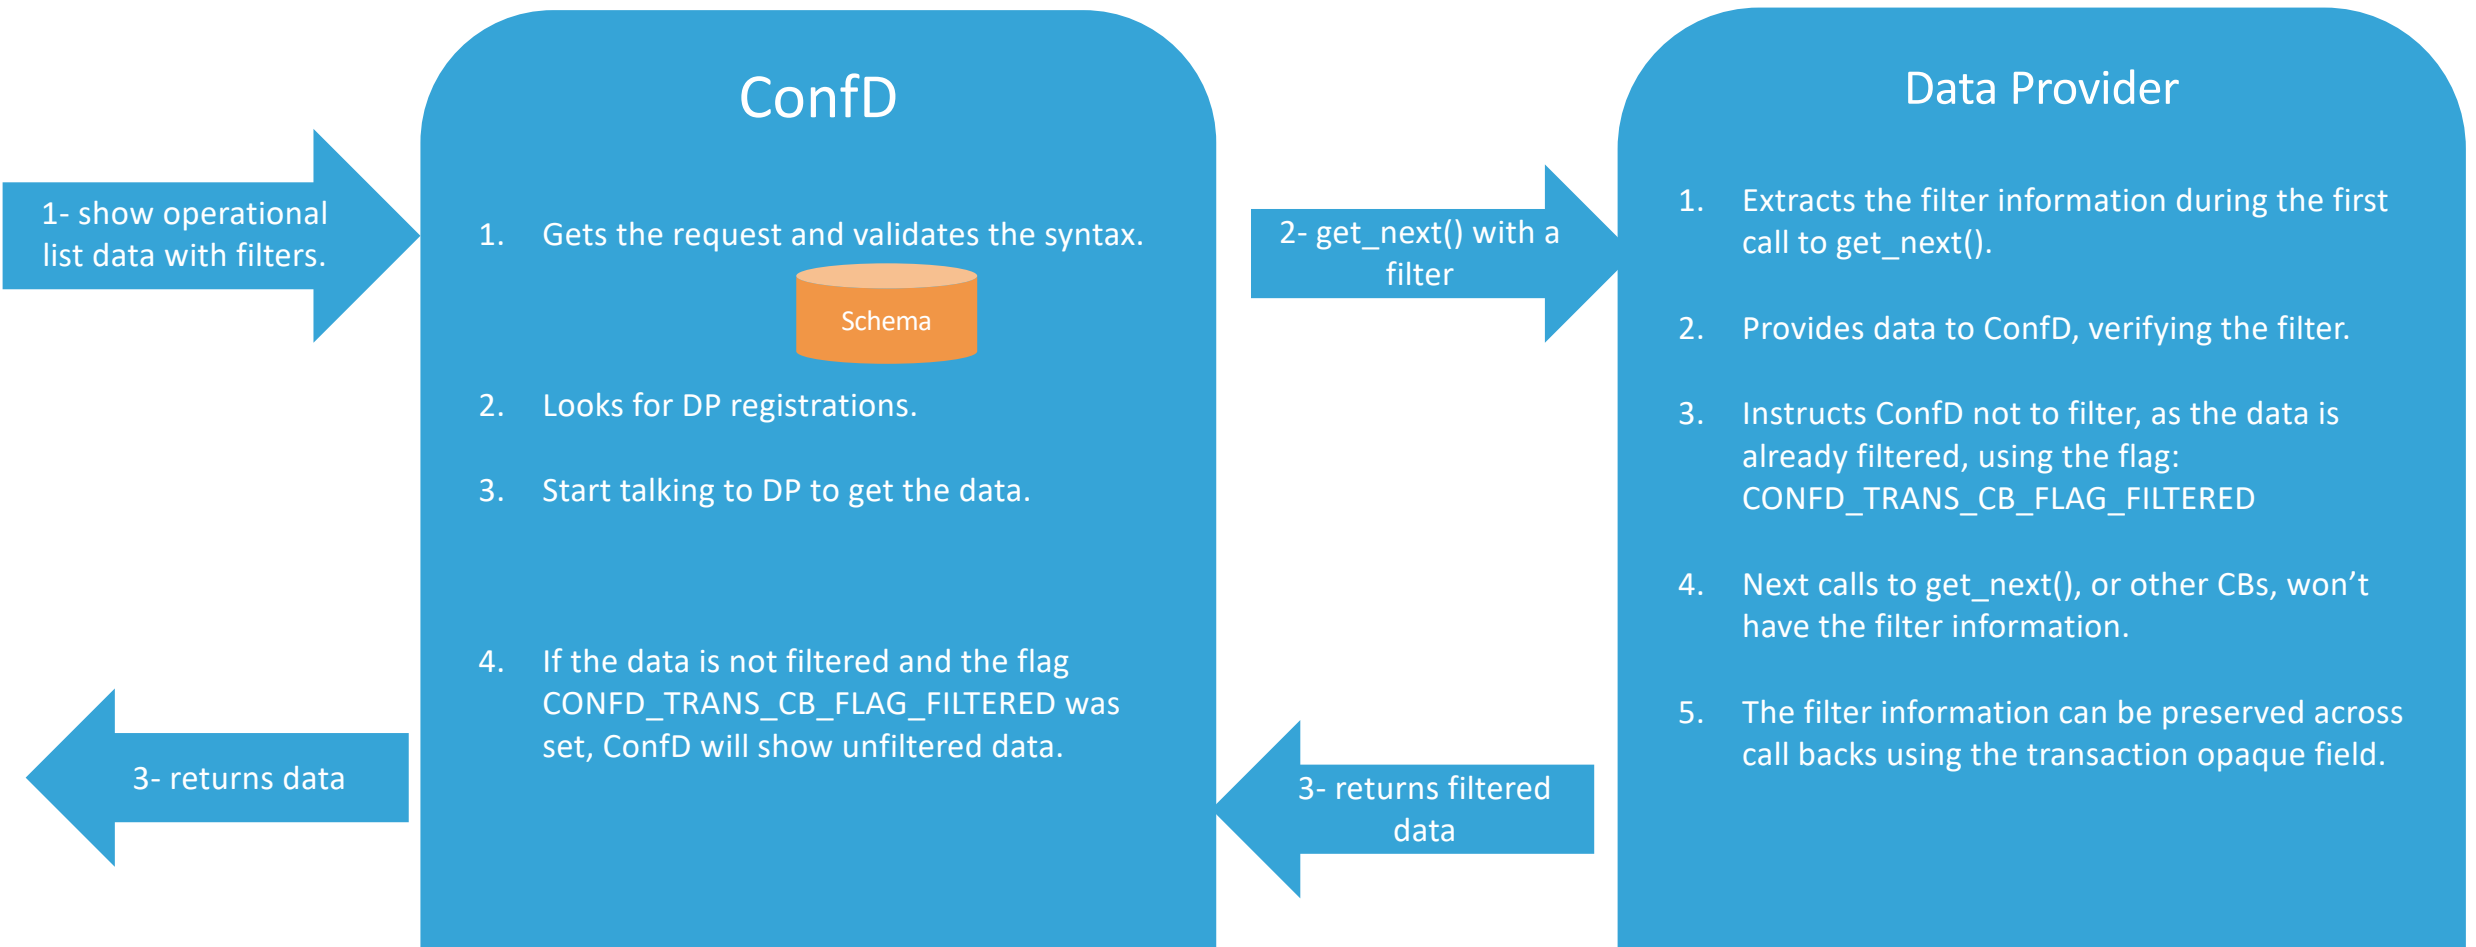
Validation

```
list foo {  
    key name;  
    leaf name { type string; }  
    leaf value { type int32; }  
}
```

```
list bar {  
    key name;  
    leaf name { type string; }  
    leaf fooref {  
        type leafref {  
            path "/foo/value";  
        }  
    }  
}
```

Delete of “/foo” with “value” 42 will get_next through the entire list bar to see if any list entry has a leaf “fooref” with value 42.

Flow



Filter syntax

```
<expr> = <expr> 'or' <expr>  
        / <expr> 'and' <expr>  
        / 'not' <expr>  
        / 'cmp' <op> <node> <value>  
        / 'exec' <func> <node> <value>  
        / 'exists' <node>
```

```
<op> = 'eq' / 'gt' / 'lt' / 'gte' / 'lte' / 'neq'
```

```
<func> = 're-match' / 'starts-with'  
        / 'derived-from' / 'derived-from-or-self'
```

```
<node> = "array of tags (points to a leaf or leaf-list)"
```

Supported Callbacks

Provided in next() callbacks:

```
get_next()  
get_next_object()  
find_next()  
find_next_object()
```

Get filter using the function (at the first invocation of the callback)

```
confd_data_get_list_filter()
```

Free the filter when done using:

```
confd_free_list_filter()
```


Performance

Operation	Regular	With filters	Ratio
Validate leaf-ref list of 10,000	22 s	1.5 s	14.7